

Введение в обработку естественного языка

Урок 10. Машинный перевод. Модель seq2seq и механизм внимания

Практическое задание

Домашнее задание к уроку 10

Задание

Разобраться с моделькой перевода как она устроена

запустить для перевода с русского на английский (при желании можно взять другие пары языков) два варианта с вниманием и без внимания оценить качество насколько корректно переводит (для теста отобрать примеры с увеличением длины текста) (так как оценка визуальная достаточно 20-ти примеров в тестовой выборке)

Выполнил *Соковнин ИЛ*

Neural machine translation

```
B [8]: import tensorflow as tf

import matplotlib.pyplot as plt
import matplotlib.ticker as ticker
from sklearn.model_selection import train_test_split

import unicodedata
import re
import numpy as np
import os
import io
import time
```

Download and prepare the dataset

Загрузка и установка датасета ¶

We'll use a language dataset provided by <http://www.manythings.org/anki/> (<http://www.manythings.org/anki/>).

```
B [2]: !wget http://www.manythings.org/anki/rus-eng.zip

--2022-06-25 15:32:24--  http://www.manythings.org/anki/rus-eng.zip (http://www.manythings.org/anki/rus-eng.zip)
Resolving www.manythings.org (www.manythings.org)... 104.21.92.44, 172.67.186.54, 2606:4700:3030::6815:5c2c, ...
Connecting to www.manythings.org (www.manythings.org)|104.21.92.44|:80... connected.
HTTP request sent, awaiting response... 200 OK
Length: 14819554 (14M) [application/zip]
Saving to: 'rus-eng.zip'

rus-eng.zip          100%[=====>]  14.13M  67.5MB/s   in 0.2s

2022-06-25 15:32:24 (67.5 MB/s) - 'rus-eng.zip' saved [14819554/14819554]
```

```
B [3]: !mkdir rus-eng
!unzip rus-eng.zip -d rus-eng/

Archive:  rus-eng.zip
  inflating: rus-eng/rus.txt
  inflating: rus-eng/_about.txt
```

```
B [4]: !ls /content/rus-eng/ -lah

total 71M
drwxr-xr-x 2 root root 4.0K Jun 25 15:32 .
drwxr-xr-x 1 root root 4.0K Jun 25 15:32 ..
-rw-r--r-- 1 root root 1.5K May  2 01:29 _about.txt
-rw-r--r-- 1 root root  71M May  2 01:29 rus.txt
```

```
B [5]: # Download the file
path_to_file = "/content/rus-eng/rus.txt"
```

```
B [6]: def preprocess_sentence(w):

    # strip() - возвращает копию строки с удаленными начальными и конечными символами (https://pythonstart.ru/string/strip)
    w = w.lower().strip()

    # creating a space between a word and the punctuation following it
    # eg: "he is a boy." => "he is a boy ."
    # Reference:- https://stackoverflow.com/questions/3645931/python-padding-punctuation-with-white-spaces-keeping-punctuation
    w = re.sub(r"([?.!,,])", r" \1 ", w)
    w = re.sub(r'[" "]+', " ", w)

    # replacing everything with space except (a-z, A-Z, ".", "?", "!", ",")
    w = re.sub(r"[^a-zA-Za-яA-Я?.!,'"]+", " ", w)

    w = w.strip()

    # adding a start and an end token to the sentence
    # so that the model know when to start and stop predicting.
    w = '<start> ' + w + ' <end>'

    return w
```

```
B [9]: preprocess_sentence("I can't go.")
```

```
Out[9]: "<start> i can't go . <end>"
```

```
B [10]: # 1. Remove the accents (Убираем акценты)
# 2. Clean the sentences (Очистите предложения)
# 3. Return word pairs in the format: [ENG, RU] (Вернуть пары слов в формате: [ENG, RU])
def create_dataset(path, num_examples):
    lines = io.open(path, encoding='UTF-8').read().strip().split('\n')

    word_pairs = [[preprocess_sentence(w) for w in l.split('\t')[:2]] for l in lines[:num_examples]]

    return zip(*word_pairs)
```

```
B [11]: lines = io.open(path_to_file, encoding='UTF-8').read().strip().split('\n')
num_examples = 1

for l in lines[:num_examples]:
    for w in l.split('\t')[:2]:
        print(preprocess_sentence(w))
```

```
<start> go . <end>
<start> mapш ! <end>
```

```
B [12]: en, ru = create_dataset(path_to_file, None)
print(en[0])
print(ru[0])
```

```
<start> go . <end>
<start> mapш ! <end>
```

```
B [13]: print(en[-1])
print(ru[-1])
```

```
<start> doubtless there exists in this world precisely the right woman for any given man to marry and vice versa but wh
en you consider that a human being has the opportunity of being acquainted with only a few hundred people , and out of
the few hundred that there are but a dozen or less whom he knows intimately , and out of the dozen , one or two friends
at most , it will easily be seen , when we remember the number of millions who inhabit this world , that probably , sin
ce the earth was created , the right man has never yet met the right woman . <end>
<start> несомненно , для каждого мужчины в этом мире где то есть подходящая женщина , которая может стать ему женой , о
братное верно и для женщин . но если учесть , что у человека может быть максимум несколько сотен знакомых , из которых
лишь дюжина , а то и меньше , тех , кого он знает близко , а из этой дюжины у него один или от силы два друга , то можн
о легко увидеть , что с уч том миллионов живущих на земле людей , ни один подходящий мужчина , возможно , ещ не встрети
л подходящую женщину . <end>
```

```
B [14]: def tokenize(lang):
        """
        Токенизация текста

        https://www.tensorflow.org/api_docs/python/tf/keras/preprocessing/text/Tokenizer
        https://www.tensorflow.org/api_docs/python/tf/keras/utils/pad_sequences

        """

        # Служебный класс для токенизации текста.
        # Этот класс позволяет векторизовать текстовый корпус.
        lang_tokenizer = tf.keras.preprocessing.text.Tokenizer(filters='')

        lang_tokenizer.fit_on_texts(lang) # Обновляет внутренний словарь на основе списка текстов.

        tensor = lang_tokenizer.texts_to_sequences(lang) # Transforms each text in texts to a sequence of integers

        tensor = tf.keras.preprocessing.sequence.pad_sequences(tensor,
                                                                padding='post') # Pads sequences to the same length (post - after)

        return tensor, lang_tokenizer
```

```
B [15]: def load_dataset(path, num_examples=None):
        # creating cleaned input, output pairs

        targ_lang, inp_lang = create_dataset(path, num_examples)

        input_tensor, inp_lang_tokenizer = tokenize(inp_lang)
        target_tensor, targ_lang_tokenizer = tokenize(targ_lang)

        return input_tensor, target_tensor, inp_lang_tokenizer, targ_lang_tokenizer
```

Limit the size of the dataset to experiment faster (optional)

Ограничим размер датасета для ускорения эксперимента (опционально)

```
B [16]: len(en), len(ru)
```

```
Out[16]: (444587, 444587)
```

```
B [17]: # Try experimenting with the size of that dataset
num_examples = 100000
input_tensor, target_tensor, inp_lang, targ_lang = load_dataset(path_to_file, num_examples)


# Calculate max_length of the target tensors
max_length_targ, max_length_inp = target_tensor.shape[1], input_tensor.shape[1]
```

```
B [18]: max_length_targ, max_length_inp
```

```
Out[18]: (11, 15)
```

```
B [19]: # Creating training and validation sets using an 80-20 split
input_tensor_train, input_tensor_val, target_tensor_train, target_tensor_val = train_test_split(input_tensor, target_tensor,
                                                                                               test_size=0.2,
                                                                                               random_state=42)

# Show Length
print(len(input_tensor_train), len(target_tensor_train), len(input_tensor_val), len(target_tensor_val))
```



```
80000 80000 20000 20000
```

```
B [20]: def convert(lang, tensor):
        for t in tensor:
            if t!=0:
                print ("%d ----> %s" % (t, lang.index_word[t]))
```

```
B [21]: print ("Input Language; index to word mapping")
convert(inp_lang, input_tensor_train[0])
print ()
print ("Target Language; index to word mapping")
convert(targ_lang, target_tensor_train[0])
```

```
Input Language; index to word mapping
1 ----> <start>
30 ----> как
210 ----> ЭТИМ
10386 ----> пользуются
5 ----> ?
2 ----> <end>
```

```
Target Language; index to word mapping
1 ----> <start>
49 ----> how
15 ----> do
4 ----> i
271 ----> use
19 ----> this
6 ----> ?
2 ----> <end>
```

Create a tf.data dataset

```
B [22]: BUFFER_SIZE = len(input_tensor_train)
BATCH_SIZE = 64
steps_per_epoch = len(input_tensor_train)//BATCH_SIZE
embedding_dim = 256
units = 1024
vocab_inp_size = len(inp_lang.word_index)+1
vocab_tar_size = len(targ_lang.word_index)+1

dataset = tf.data.Dataset.from_tensor_slices((input_tensor_train, target_tensor_train)).shuffle(BUFFER_SIZE)
dataset = dataset.batch(BATCH_SIZE, drop_remainder=True)
```

```
B [23]: example_input_batch, example_target_batch = next(iter(dataset))
example_input_batch.shape, example_target_batch.shape
```

```
Out[23]: (TensorShape([64, 15]), TensorShape([64, 11]))
```

```
B [26]: class Encoder(tf.keras.Model):
def __init__(self, vocab_size, embedding_dim, enc_units, batch_sz):
    super(Encoder, self).__init__()
    self.batch_sz = batch_sz
    self.enc_units = enc_units
    self.embedding = tf.keras.layers.Embedding(vocab_size, embedding_dim)
    self.gru = tf.keras.layers.GRU(self.enc_units,
                                   return_sequences=False,
                                   return_state=True,
                                   recurrent_initializer='glorot_uniform')

def call(self, x, hidden):
    x = self.embedding(x)
    output, state = self.gru(x, initial_state = hidden)
    return state

def initialize_hidden_state(self):
    return tf.zeros((self.batch_sz, self.enc_units))
```

Neural machine translation

```
B [27]: encoder = Encoder(vocab_inp_size, embedding_dim, units, BATCH_SIZE)

# sample input
sample_hidden = encoder.initialize_hidden_state()
sample_hidden = encoder(example_input_batch, sample_hidden)
# print ('Encoder output shape: (batch size, sequence length, units) {}'.format(sample_output.shape))
print ('Encoder Hidden state shape: (batch size, units) {}'.format(sample_hidden.shape))
```

```
Encoder Hidden state shape: (batch size, units) (64, 1024)
```

```
B [28]: class Decoder(tf.keras.Model):
def __init__(self, vocab_size, embedding_dim, dec_units, batch_sz):
    super(Decoder, self).__init__()
    self.batch_sz = batch_sz
    self.dec_units = dec_units
    self.embedding = tf.keras.layers.Embedding(vocab_size, embedding_dim)
    self.gru = tf.keras.layers.GRU(self.dec_units,
                                   return_sequences=True,
                                   return_state=True,
                                   recurrent_initializer='glorot_uniform')

    self.fc = tf.keras.layers.Dense(vocab_size)

def call(self, x, hidden):
    # enc_output shape == (batch_size, max_length, hidden_size)

    # x shape after passing through embedding == (batch_size, 1, embedding_dim)
    x = self.embedding(x)

    # x shape after concatenation == (batch_size, 1, embedding_dim + hidden_size)

    # passing the concatenated vector to the GRU
    output, state = self.gru(x, initial_state=hidden)

    # output shape == (batch_size * 1, hidden_size)
    output = tf.reshape(output, (-1, output.shape[2]))

    # output shape == (batch_size, vocab)
    x = self.fc(output)

    return x, state
```

```
B [29]: decoder = Decoder(vocab_tar_size, embedding_dim, units, BATCH_SIZE)

decoder_sample_x, decoder_sample_h = decoder(tf.random.uniform((BATCH_SIZE, 1)),
                                             sample_hidden)
```

```
B [30]: decoder_sample_x.shape
```

```
Out[30]: TensorShape([64, 7334])
```

```
B [31]: decoder_sample_h.shape
```

```
Out[31]: TensorShape([64, 1024])
```

Define the optimizer and the loss function

```
B [32]: optimizer = tf.keras.optimizers.Adam()

loss_object = tf.keras.losses.SparseCategoricalCrossentropy(
    from_logits=True, reduction='none')

def loss_function(real, pred):
    mask = tf.math.logical_not(tf.math.equal(real, 0))
    loss_ = loss_object(real, pred)

    mask = tf.cast(mask, dtype=loss_.dtype)
    loss_ *= mask

    return tf.reduce_mean(loss_)
```

Checkpoints (Object-based saving)

Контрольные точки (Object-based saving)

```
B [33]: checkpoint_dir = './training_checkpoints'

checkpoint_prefix = os.path.join(checkpoint_dir, "ckpt")

checkpoint = tf.train.Checkpoint(optimizer=optimizer,
                                 encoder=encoder,
                                 decoder=decoder)
```

Обучение модели

```

B [34]: @tf.function
def train_step(inp, targ, enc_hidden):
    loss = 0

    with tf.GradientTape() as tape:
        enc_hidden = encoder(inp, enc_hidden)

        dec_hidden = enc_hidden

        dec_input = tf.expand_dims([targ_lang.word_index['<start>']] * BATCH_SIZE, 1)

        # Teacher forcing - feeding the target as the next input
        for t in range(1, targ.shape[1]):
            # passing enc_output to the decoder
            predictions, dec_hidden = decoder(dec_input, dec_hidden)

            loss += loss_function(targ[:, t], predictions)

            # using teacher forcing
            dec_input = tf.expand_dims(targ[:, t], 1)

        batch_loss = (loss / int(targ.shape[1]))

        variables = encoder.trainable_variables + decoder.trainable_variables

        gradients = tape.gradient(loss, variables)

        optimizer.apply_gradients(zip(gradients, variables))

    return batch_loss

```

```

B [36]: EPOCHS = 20

for epoch in range(EPOCHS):
    start = time.time()

    enc_hidden = encoder.initialize_hidden_state()
    total_loss = 0

    for (batch, (inp, targ)) in enumerate(dataset.take(steps_per_epoch)):
        batch_loss = train_step(inp, targ, enc_hidden)
        total_loss += batch_loss

        if batch % 100 == 0:
            print('Epoch {} Batch {} Loss {:.4f}'.format(epoch + 1,
                                                            batch,
                                                            batch_loss.numpy()))

    # saving (checkpoint) the model every 2 epochs
    if (epoch + 1) % 2 == 0:
        checkpoint.save(file_prefix = checkpoint_prefix)

    print('Epoch {} Loss {:.4f}'.format(epoch + 1,
                                         total_loss / steps_per_epoch))
    print('Time taken for 1 epoch {} sec\n'.format(time.time() - start))

```

```

Epoch 1 Batch 0 Loss 4.5258
Epoch 1 Batch 100 Loss 2.0305
Epoch 1 Batch 200 Loss 1.7621
Epoch 1 Batch 300 Loss 1.5777
Epoch 1 Batch 400 Loss 1.5404
Epoch 1 Batch 500 Loss 1.4149
Epoch 1 Batch 600 Loss 1.3890
Epoch 1 Batch 700 Loss 1.2275
Epoch 1 Batch 800 Loss 1.2151
Epoch 1 Batch 900 Loss 1.0821
Epoch 1 Batch 1000 Loss 1.1432
Epoch 1 Batch 1100 Loss 1.1177
Epoch 1 Batch 1200 Loss 1.0775
Epoch 1 Loss 1.4569
Time taken for 1 epoch 61.77876329421997 sec

```

```

Epoch 2 Batch 0 Loss 0.8406
Epoch 2 Batch 100 Loss 0.9508
Epoch 2 Batch 200 Loss 0.8073
Epoch 2 Batch 300 Loss 0.8021

```

Translate

- The evaluate function is similar to the training loop, except we don't use *teacher forcing* here. The input to the decoder at each time step is its previous predictions along with the hidden state and the encoder output.
- Stop predicting when the model predicts the *end token*.
- And store the *attention weights for every time step*.

Note: The encoder output is calculated only once for one input.

```
B [37]: def evaluate(sentence):
    attention_plot = np.zeros((max_length_targ, max_length_inp))

    sentence = preprocess_sentence(sentence)

    inputs = [inp_lang.word_index[i] for i in sentence.split(' ')]
    inputs = tf.keras.preprocessing.sequence.pad_sequences([inputs],
                                                            maxlen=max_length_inp,
                                                            padding='post')

    inputs = tf.convert_to_tensor(inputs)

    result = ''

    hidden = [tf.zeros((1, units))]
    enc_hidden = encoder(inputs, hidden)

    dec_hidden = enc_hidden
    dec_input = tf.expand_dims([targ_lang.word_index['<start>']], 0)

    for t in range(max_length_targ):
        predictions, dec_hidden = decoder(dec_input, dec_hidden)

        # storing the attention weights to plot later on
        predicted_id = tf.argmax(predictions[0]).numpy()
        result += targ_lang.index_word[predicted_id] + ' '

        if targ_lang.index_word[predicted_id] == '<end>':
            return result, sentence

        # the predicted ID is fed back into the model
        dec_input = tf.expand_dims([predicted_id], 0)

    return result, sentence
```

```
B [38]: def translate(sentence):
    result, sentence = evaluate(sentence)

    print('Input: %s' % (sentence))
    print('Predicted translation: {}'.format(result))
```

Restore the latest checkpoint and test

```
B [39]: # restoring the latest checkpoint in checkpoint_dir
checkpoint.restore(tf.train.latest_checkpoint(checkpoint_dir))
```

Out[39]: <tensorflow.python.training.tracking.util.CheckpointLoadStatus at 0x7f06e8555110>

```
B [40]: translate('Здесь хорошо.')
```

Input: <start> здесь хорошо . <end>
Predicted translation: it's fine here . <end>

```
B [41]: translate('Отлично, поехали.')
```

Input: <start> отлично , поехали . <end>
Predicted translation: ok , let's go . <end>

```
B [42]: translate(u'Вы еще дома?')
```

Input: <start> вы еще дома ? <end>
Predicted translation: are you still at home ? <end>

```
B [43]: translate(u'Это слишком дорого для меня.?')
```

Input: <start> это слишком дорого для меня . ? <end>
Predicted translation: that's too much for me . <end>

```
B [44]: translate(u'Попробуй сделать это.')
```

Input: <start> попробуй сделать это . <end>
Predicted translation: try to do that . <end>

```
B [45]: translate(u'Я люблю, когда идет снег.')
```

Input: <start> я люблю , когда идет снег . <end>
Predicted translation: i like it dark . <end>

B [46]: `translate(u'Я никогда такого не делаю.')`

Input: <start> я никогда такого не делаю . <end>
Predicted translation: i never do that . <end>

B [47]: `translate('А счастье было так возможно, так близко!.')`

Input: <start> а счастье было так возможно , так близко ! . <end>
Predicted translation: it was cheap . <end>

B [48]: `translate('Интересно, а если написать длинное предложение и попробовать его перевести, какой результат мы увидим?')`

Input: <start> интересно , а если написать длинное предложение и попробовать его перевести , какой результат мы увидим ? <end>
Predicted translation: do you also like jazz ? <end>

Вывод:

При увеличении длины текста, качество перевода резко падает. Перевод последнего предложения соответствует последнему предложению.

Neural machine translation with attention

```
B [49]: class Encoder(tf.keras.Model):
def __init__(self, vocab_size, embedding_dim, enc_units, batch_sz):
    super(Encoder, self).__init__()
    self.batch_sz = batch_sz
    self.enc_units = enc_units
    self.embedding = tf.keras.layers.Embedding(vocab_size, embedding_dim)
    self.gru = tf.keras.layers.GRU(self.enc_units,
                                   return_sequences=True,
                                   return_state=True,
                                   recurrent_initializer='glorot_uniform')

def call(self, x, hidden):
    x = self.embedding(x)
    output, state = self.gru(x, initial_state = hidden)
    return output, state

def initialize_hidden_state(self):
    return tf.zeros((self.batch_sz, self.enc_units))
```

```
B [50]: encoder = Encoder(vocab_inp_size, embedding_dim, units, BATCH_SIZE)

# sample input
sample_hidden = encoder.initialize_hidden_state()
sample_output, sample_hidden = encoder(example_input_batch, sample_hidden)
print ('Encoder output shape: (batch size, sequence length, units) {}'.format(sample_output.shape))
print ('Encoder Hidden state shape: (batch size, units) {}'.format(sample_hidden.shape))
```

Encoder output shape: (batch size, sequence length, units) (64, 15, 1024)
Encoder Hidden state shape: (batch size, units) (64, 1024)


```

B [51]: class BahdanauAttention(tf.keras.layers.Layer):
    def __init__(self, units):
        super(BahdanauAttention, self).__init__()
        self.W1 = tf.keras.layers.Dense(units)
        self.W2 = tf.keras.layers.Dense(units)
        self.V = tf.keras.layers.Dense(1)

    def call(self, query, values):
        # query hidden state shape == (batch_size, hidden_size)
        # query_with_time_axis shape == (batch_size, 1, hidden_size)
        # values shape == (batch_size, max_len, hidden_size)
        # we are doing this to broadcast addition along the time axis to calculate the score
        query_with_time_axis = tf.expand_dims(query, 1)

        # score shape == (batch_size, max_length, 1)
        # we get 1 at the last axis because we are applying score to self.V
        # the shape of the tensor before applying self.V is (batch_size, max_length, units)
        score = self.V(tf.nn.tanh(
            self.W1(query_with_time_axis) + self.W2(values)))

        # attention_weights shape == (batch_size, max_length, 1)
        attention_weights = tf.nn.softmax(score, axis=1)

        # context_vector shape after sum == (batch_size, hidden_size)
        context_vector = attention_weights * values
        context_vector = tf.reduce_sum(context_vector, axis=1)

        return context_vector, attention_weights

```

```

B [52]: attention_layer = BahdanauAttention(10)
attention_result, attention_weights = attention_layer(sample_hidden, sample_output)

print("Attention result shape: (batch size, units) {}".format(attention_result.shape))
print("Attention weights shape: (batch_size, sequence_length, 1) {}".format(attention_weights.shape))

```

Attention result shape: (batch size, units) (64, 1024)
 Attention weights shape: (batch_size, sequence_length, 1) (64, 15, 1)

```

B [53]: class Decoder(tf.keras.Model):
    def __init__(self, vocab_size, embedding_dim, dec_units, batch_sz):
        super(Decoder, self).__init__()
        self.batch_sz = batch_sz
        self.dec_units = dec_units
        self.embedding = tf.keras.layers.Embedding(vocab_size, embedding_dim)
        self.gru = tf.keras.layers.GRU(self.dec_units,
                                         return_sequences=True,
                                         return_state=True,
                                         recurrent_initializer='glorot_uniform')
        self.fc = tf.keras.layers.Dense(vocab_size)

        # used for attention
        self.attention = BahdanauAttention(self.dec_units)

    def call(self, x, hidden, enc_output):
        # enc_output shape == (batch_size, max_length, hidden_size)
        context_vector, attention_weights = self.attention(hidden, enc_output)

        # x shape after passing through embedding == (batch_size, 1, embedding_dim)
        x = self.embedding(x)

        # x shape after concatenation == (batch_size, 1, embedding_dim + hidden_size)
        x = tf.concat([tf.expand_dims(context_vector, 1), x], axis=-1)

        # passing the concatenated vector to the GRU
        output, state = self.gru(x)

        # output shape == (batch_size * 1, hidden_size)
        output = tf.reshape(output, (-1, output.shape[2]))

        # output shape == (batch_size, vocab)
        x = self.fc(output)

        return x, state, attention_weights

```

```

B [54]: decoder = Decoder(vocab_tar_size, embedding_dim, units, BATCH_SIZE)

sample_decoder_output, _, _ = decoder(tf.random.uniform((BATCH_SIZE, 1)),
                                       sample_hidden, sample_output)

print('Decoder output shape: (batch_size, vocab size) {}'.format(sample_decoder_output.shape))

```

Decoder output shape: (batch_size, vocab size) (64, 7334)

Define the optimizer and the loss function

```
B [55]: optimizer = tf.keras.optimizers.Adam()
loss_object = tf.keras.losses.SparseCategoricalCrossentropy(
    from_logits=True, reduction='none')

def loss_function(real, pred):
    mask = tf.math.logical_not(tf.math.equal(real, 0))
    loss_ = loss_object(real, pred)

    mask = tf.cast(mask, dtype=loss_.dtype)
    loss_ *= mask

    return tf.reduce_mean(loss_)
```

Checkpoints (Object-based saving)

```
B [56]: checkpoint_dir = './training_attention_checkpoints'
checkpoint_prefix = os.path.join(checkpoint_dir, "ckpt")
checkpoint = tf.train.Checkpoint(optimizer=optimizer,
                                encoder=encoder,
                                decoder=decoder)
```

```
B [58]: @tf.function
def train_step(inp, targ, enc_hidden):
    loss = 0

    with tf.GradientTape() as tape:
        enc_output, enc_hidden = encoder(inp, enc_hidden)

        dec_hidden = enc_hidden

        dec_input = tf.expand_dims([targ_lang.word_index['<start>']] * BATCH_SIZE, 1)

        # Teacher forcing - feeding the target as the next input
        for t in range(1, targ.shape[1]):
            # passing enc_output to the decoder
            predictions, dec_hidden, _ = decoder(dec_input, dec_hidden, enc_output)

            loss += loss_function(targ[:, t], predictions)

            # using teacher forcing
            dec_input = tf.expand_dims(targ[:, t], 1)

    batch_loss = (loss / int(targ.shape[1]))

    variables = encoder.trainable_variables + decoder.trainable_variables

    gradients = tape.gradient(loss, variables)

    optimizer.apply_gradients(zip(gradients, variables))

    return batch_loss
```

B [59]: EPOCHS = 20

```
for epoch in range(EPOCHS):
    start = time.time()

    enc_hidden = encoder.initialize_hidden_state()
    total_loss = 0

    for (batch, (inp, targ)) in enumerate(dataset.take(steps_per_epoch)):
        batch_loss = train_step(inp, targ, enc_hidden)
        total_loss += batch_loss

        if batch % 100 == 0:
            print('Epoch {} Batch {} Loss {:.4f}'.format(epoch + 1,
                                                            batch,
                                                            batch_loss.numpy()))

    # saving (checkpoint) the model every 2 epochs
    if (epoch + 1) % 2 == 0:
        checkpoint.save(file_prefix = checkpoint_prefix)

    print('Epoch {} Loss {:.4f}'.format(epoch + 1,
                                         total_loss / steps_per_epoch))
    print('Time taken for 1 epoch {} sec\n'.format(time.time() - start))
```

```
Epoch 1 Batch 0 Loss 4.8045
Epoch 1 Batch 100 Loss 2.0884
Epoch 1 Batch 200 Loss 1.8418
Epoch 1 Batch 300 Loss 1.7152
Epoch 1 Batch 400 Loss 1.5148
Epoch 1 Batch 500 Loss 1.5344
Epoch 1 Batch 600 Loss 1.5292
Epoch 1 Batch 700 Loss 1.2826
Epoch 1 Batch 800 Loss 1.2203
Epoch 1 Batch 900 Loss 1.1733
Epoch 1 Batch 1000 Loss 1.1406
Epoch 1 Batch 1100 Loss 1.1039
Epoch 1 Batch 1200 Loss 1.0212
Epoch 1 Loss 1.4665
Time taken for 1 epoch 112.27705335617065 sec
```

```
Epoch 2 Batch 0 Loss 0.7619
Epoch 2 Batch 100 Loss 0.9176
Epoch 2 Batch 200 Loss 0.8496
Epoch 2 Batch 300 Loss 0.7600
```

Translate

- The evaluate function is similar to the training loop, except we don't use *teacher forcing* here. The input to the decoder at each time step is its previous predictions along with the hidden state and the encoder output.
- Stop predicting when the model predicts the *end token*.
- And store the *attention weights for every time step*.

Note: The encoder output is calculated only once for one input.

```

B [60]: def evaluate(sentence):
    attention_plot = np.zeros((max_length_targ, max_length_inp))

    sentence = preprocess_sentence(sentence)

    inputs = [inp_lang.word_index[i] for i in sentence.split(' ')]
    inputs = tf.keras.preprocessing.sequence.pad_sequences([inputs],
                                                            maxlen=max_length_inp,
                                                            padding='post')

    inputs = tf.convert_to_tensor(inputs)

    result = ''

    hidden = [tf.zeros((1, units))]
    enc_out, enc_hidden = encoder(inputs, hidden)

    dec_hidden = enc_hidden
    dec_input = tf.expand_dims([targ_lang.word_index['<start>']], 0)

    for t in range(max_length_targ):
        predictions, dec_hidden, attention_weights = decoder(dec_input,
                                                            dec_hidden,
                                                            enc_out)

        # storing the attention weights to plot later on
        attention_weights = tf.reshape(attention_weights, (-1, ))
        attention_plot[t] = attention_weights.numpy()

        predicted_id = tf.argmax(predictions[0]).numpy()

        result += targ_lang.index_word[predicted_id] + ' '

        if targ_lang.index_word[predicted_id] == '<end>':
            return result, sentence, attention_plot

        # the predicted ID is fed back into the model
        dec_input = tf.expand_dims([predicted_id], 0)

    return result, sentence, attention_plot

```

```

B [61]: # function for plotting the attention weights
def plot_attention(attention, sentence, predicted_sentence):
    fig = plt.figure(figsize=(10,10))
    ax = fig.add_subplot(1, 1, 1)
    ax.matshow(attention, cmap='viridis')

    fontdict = {'fontsize': 14}

    ax.set_xticklabels([''] + sentence, fontdict=fontdict, rotation=90)
    ax.set_yticklabels([''] + predicted_sentence, fontdict=fontdict)

    ax.xaxis.set_major_locator(ticker.MultipleLocator(1))
    ax.yaxis.set_major_locator(ticker.MultipleLocator(1))

    plt.show()

```

```

B [62]: def translate(sentence):
    result, sentence, attention_plot = evaluate(sentence)

    print('Input: %s' % (sentence))
    print('Predicted translation: {}'.format(result))

    attention_plot = attention_plot[:len(result.split(' ')), :len(sentence.split(' '))]
    plot_attention(attention_plot, sentence.split(' '), result.split(' '))

```

Restore the latest checkpoint and test

```


B [63]: # restoring the latest checkpoint in checkpoint_dir
checkpoint.restore(tf.train.latest_checkpoint(checkpoint_dir))

```

Out[63]: <tensorflow.python.training.tracking.util.CheckpointLoadStatus at 0x7f06f3b2fa50>


```
B [64]: translate('Здесь хорошо.')
```

Input: <start> здесь хорошо . <end>
Predicted translation: it's fine here . <end>




```
B [65]: translate('Я не смогу поехать.')
```

Input: <start> я не смогу поехать . <end>
Predicted translation: i can't go . <end>




```
B [66]: translate(u'Вы еще дома?')
```

Input: <start> вы еще дома ? <end>
Predicted translation: are you still home ? <end>




B [67]: `translate(u'Вы все еще дома?')`

Input: <start> вы все еще дома ? <end>
Predicted translation: are you still at home ? <end>



```
B [68]: translate(u'Попробуй сделать это.')
```


Input: <start> попробуй сделать это . <end>
Predicted translation: try doing it fixed . <end>



```
B [69]: translate(u'Попробуй сделать это.')

```


Input: <start> попробуй сделать это . <end>
Predicted translation: try doing it fixed . <end>



```
translate(u'Я не люблю, когда идет снег.')
```

Input: <start> я не люблю , когда идет снег . <end>


Input: <start> Я не люблю , когда идет снег : <end>
Predicted translation: i don't like snow . <end>



```
translate(u'Я никогда такого не делаю.')
```


Input: <start> я никогда такого не делаю . <end>

Input: <stat> я никогда такого не делаю : <end>
Predicted translation: i never do that . <end>



```
В [72]: translate(u'Интересно, а если написать длинное предложение и попробовать его перевести, какой результат мы увидим?')
```

Input: <start> интересно , а если написать длинное предложение и попробовать его перевести , какой результат мы увидим ? <end>
Predicted translation: i have to act . <end>



Вывод:
При увеличении длины текста, качество перевода резко падает. Перевод не соответствует последнему предложению.

```
В [ ]:
```